

EDUCATION

University of Southern California	2025 - 2027
M.S. in Analytics, Los Angeles, CA	
Vellore Institute of Technology	2021 - 2025
B.Tech in Computer Science & Engineering	

TECHNICAL SKILLS

AI/Automation Tools	ChatGPT, GitHub Copilot, Microsoft Copilot, Vertex AI, Power Automate, Prompt Engineering
Workflow/Documentation	Process Automation, Structured Research Workflows, Source-to-Target Mapping, Metadata Strategy
Languages/Scripting	Python (Pandas, NumPy, Scikit-learn), SQL, Git, Java
Reporting/Bi Tools	Tableau (DataCamp Certified), Power BI, dbt, Snowflake
Data Infrastructure	GCP, AWS, DynamoDB, AWS Kinesis, Lambda, Terraform (IaC), Neo4j (Graph DB)
Certifications	AWS Cloud Practitioner (Amazon), Tableau & SQL (DataCamp)

WORK EXPERIENCE

Integrated Learning Systems India Pvt. Ltd	Jan 2025 - Apr 2025
Data Engineering Intern - ML & Security Analytics	
- Identified fraud-like behavioral patterns in network traffic by constructing a parameterized Vertex AI + KFP pipeline automating ingestion, StandardScaler transformation, and Logistic Regression classification, achieving 97.0% validation accuracy and F1 of 0.9767 with zero training-serving skew. - Deployed a production GCP inference runtime that dynamically reads model artifacts from GCS and replicates upstream scaler states, eliminating numeric drift across inference boundaries on live phishing and malware datasets. - Accelerated ML pipeline delivery 2× by embedding LLM-driven workflows to compile JSON KFP schemas and automate GCS artifact uploads, enforcing full reproducibility across iterative cybersecurity model retraining cycles.	
ATCreations Design School	Mar 2024 - Aug 2024
Data & UX Analytics Intern	
- Formulated a multi-page Tableau BI platform across 4 regions tracking Average Monthly Revenue and Marketing ROI, exposing a marketing efficiency gap of \$0.41 vs. \$0.62/spent and a holiday promotional lift of \$275 to \$311. - Converted qualitative client requirements into quantitative analytical frameworks across 180+ inventory records, producing UX design recommendations that improved user engagement by 16%. - Accelerated dashboard delivery by deploying AI co-pilots to generate logical syntax trees, resolve tabular join mismatches, and cross-validate calculated field dependencies across e-commerce brand analytics initiatives.	
SmartED Innovations	Aug 2023 - Nov 2023
Junior Data Engineer Intern	
- Engineered Python/Pandas ETL pipelines unpacking nested JSON payloads via <code>json_normalize</code> into clean relational dataframes with enforced UTF-8 encoding, eliminating schema mismatches across heterogeneous cybersecurity and network optimization data sources. - Constructed Snowflake SQL analytical queries using window functions and CTEs to compute running totals, streak analysis, and multi-tiered profiles, ensuring data integrity and reliability across large-scale security monitoring pipelines. - Accelerated pipeline delivery 2× by leveraging AI co-pilots to auto-generate regex cleaning functions and extraction loops; ranked top 20 out of 100 trainees for outstanding problem-solving performance.	

PROJECTS

AI-Era Identity Intelligence & Real-Time Financial Fraud Streaming Platform Python, AWS Kinesis, XGBoost	
- Architected a cloud-agnostic streaming fraud detection platform on AWS Kinesis/Lambda with a Neo4j identity graph and DynamoDB feature store, processing 35,000 events/second at p99 latency of 0.04ms, 5000× faster than standard banking compliance targets across ISO 20022, FedNow, and RTP payment rails. - Deployed a dual-layer XGBoost + deterministic rules scoring engine and leveraged generative AI co-pilots to synthesize multi-cloud Terraform IaC templates and a 54-stage pytest regression suite, achieving full integration coverage.	
Analytical Data Warehouse Architecture & Conformed Star Schema Snowflake, SQL, Kimball Modeling, dbt, Python	
- Designed a Kimball-modeled star schema in Snowflake with a central FACT_SALES table, conformed dimensions, and custom U.S. federal holiday flags, fully decoupling transactional locks from analytical reporting workloads. - Implemented additive and non-additive metrics mapped via surrogate keys and paired generative AI assistants to synthesize source-to-target mapping matrices and Snowflake DDL scripts, compressing build timelines by an estimated 2×	
Cloud-Native Healthcare Data Engineering & Predictive Pipeline GCP, Vertex AI, Kubeflow Pipelines, XGBoost	
- Engineered a decoupled KFP orchestration framework on GCP with containerized DAG components handling SMOTE oversampling and median imputation, resolving 30.9% data missingness across a 4.8:1 class-imbalanced clinical dataset without validation leakage. - Optimized an XGBoost model to F1 of 0.8493 and ROC-AUC of 0.9795 by benchmarking 5 classifiers under Vertex AI Experiments and tuning decision threshold 0.50 → 0.44 to minimize clinical false-negatives.	

LEADERSHIP & INVOLVEMENT

Head Coordinator - “BGMI Showdown” & “MELTDOWN” (National Debate & GD Competition); led end-to-end planning and execution at national scale. “TECH LITE” (National Technical Event) & “DESIGNATHON” (National UI/UX Contest), CSE Dept.